

```
!pip install streamlit
!pip install PyPDF2
!pip install langchain
!pip install -U langchain-community
!pip install faiss-cpu
!pip install openai
```

 [Show hidden output](#)

```
%%writefile app.py
import streamlit as st
from PyPDF2 import PdfReader
from langchain.text_splitter import RecursiveCharacterTextSplitter
from langchain.embeddings.openai import OpenAIEmbeddings
from langchain_community.vectorstores import FAISS
from langchain.chains.question_answering import load_qa_chain
from langchain_community.chat_models import ChatOpenAI

OPENAI_API_KEY=" your API KEY"

#upload pdf files
st.header("ChatBot")
with st.sidebar:
    st.title("File Upload")
    file=st.file_uploader("Choose a pdf file", type="pdf")

# #extracting the text from pdf files.
text = ""
if file is not None:
    pdf_reader=PdfReader(file)
    for page in pdf_reader.pages:
        text+=page.extract_text()
else:
    st.warning("Please upload a pdf file")
    ...

breaking the text into chunks
text splitter will be used from langchain.
...

text_splitter=RecursiveCharacterTextSplitter(
    separators="\n",
    chunk_size=1000,
    chunk_overlap=150, # so that the relation and the meaning of the chunks is preserved, to save incomplete sentences.
    length_function=len # breaking based on 1000 characters as length is given
)
chunks=text_splitter.split_text(text)
#-----
...

post geberaton of the chunks, we will create those embeddings and then store them in a vector store
this is like a database to store the embeddings.
...

#-----
# using openai services with langchain to do so
embeddings=OpenAIEmbeddings(openai_api_key=OPENAI_API_KEY)
# creating our vector store using FAISS - Facebook AI Semantic Search, Again Langchain will be used.
vector_store=FAISS.from_texts(chunks, embeddings) # it will take in the chunks and the embeddings.
...

# - generated embeddings with OpenAI services
# - created/initialized the FAISS vector store
# - stores the chunks and embeddings
...

#-----
...

We will proceed to Finetuning section
1. get user's question
2. do similarity search
3. output results
...

#defining our LLM:
llm = ChatOpenAI(
    open_api_key=OPENAI_API_KEY,
    temperature=0, # reducing the randomness, lower the value more specific the answer and not random
    max_tokens=1000,
    model_name="gpt-3.5-turbo"
)
user_question=st.text_input("Ask your Question")
if user_question:
    matches=vector_store.similarity_search(user_question)
    chain=load_qa_chain=load_qa_chain(llm, chain_type="stuff")
    response=chain.run(input_documents=matches, question=user_question)
    st.write(response)
```

↔ Overwriting app.py

```
!npm install localtunnel
```

```
!streamlit run app.py &> logs.txt & curl ipv4.icanhazip.com
```

```
!npx localtunnel --port (port)
```